

# AI4001/CS4063 - Fundamentals of NLP/NLP Course Project

## 1 Objective

The project is an important part of this course and makes up a significant portion of your final mark. The project gives you a chance to **develop a small AI Agent** from among the options given and **rigorously test** it for **scalability, reliability** and **performance**. Your work should be **well documented for reproducibility**.

In particular, you will be expected to become familiar with **existing work on Agentic solutions** and develop a **small system that is containerized** and can be **used as an API and tested via Postman**.

You may do your implementation within any freely-available infrastructures, toolkits and models (**Phi**, **Llama**, **StableDiffusion** etc), however, the solution should be self-contained and automatically deployable via **workflows/deployment scripts**. Your microservice should be **detached** from the **front-end that you would develop to demo the functionality** of your microservice.

## 2 Microservice Requirements

The microservice you develop is expected to:

- Have gRPC API with at least 1 core endpoint (e.g., /generate, /infer)
- Handle multiple requests concurrently (basic async/concurrency)
- Support minimal error handling (invalid input, server errors)
- Return JSON responses with status codes
- Should be thoroughly tested (**testcases to be submitted**)
- Should be demoable **through Postman + a minimal fronted using Gradio or StreamLit**,

## 3 Project Ideas

Projects should be done in groups of two or three and must solve one of the following problems:

- **Story2Audio** — Given a storyline, convert it to an engaging audio-story (like an audio book). The audio tone should match the text being read and not in monotonic voice.
- **TrendStory** — Write a program to extract current trends (for the past 24hrs, from youtube, tiktok or Google) and generate a story script in a particular theme of your choice (comedy, sarcasm, tragedy etc)
- **Text2Image** — Given a sentence or a paragraph along with context (text, context-of-text) generate an appropriate image.

- Text2Video — Given a sentence or a paragraph along with context (text, context-of-text) generate an appropriate video.
- Image2Video — Given an image and its context in text (image,context-of-image) generate an appropriate video.

As a suggestion use off-the-shelf open-weights models, GitHub Actions for CI/CD and Huggingface Spaces to deploy your model.

## 4 Submission Guidelines

The deliverables expected are:

1. Source code in a GitHub repo (also to be submitted on classroom)
2. Gradio/StreaLit frontend application + Postman collection with test cases
3. Dockerfile and/or deployment scripts
4. Documentation (README with setup, usage, architecture, model sources, limitations)
5. Testcases and performance evaluation (performance graphs like # of concurrent requests/response time)
6. Short video demo (optional but helpful)

### Evaluation Criteria:

- Functionality and Correctness: 30%
- API Design and Usability: 10%
- Code Quality and Documentation: 20%
- Testing and Fault Tolerance: 20%
- Reproducibility and Deployment: 20%

## 5 Project Presentation

Near the end of the term, you will give a demo of your project in a 3-minute presentation to the class and an evaluation panel.